

Consiva · Lead Generation Pipeline

MCA-Sourced, Consent-Gap Verified, Contact-Enriched — Executive Summary

A single, focused pipeline: official company data goes in, and a verified DPDP compliance gap plus a real, enriched contact comes out. Internal build reference (v2) — Swaran Soft.

1. The Pipeline at a Glance — 7 Stages

#	Stage	What Happens
1	MCA Registered Companies	Official data.gov.in company master list, filtered to Active status
2	Domain Resolution	Name → slug → direct HTTP probe, no search engine involved
3	Playwright Multi-Purpose Scrape	Consent-gap detection + tech stack + content signals in one page load
4	Apify LinkedIn Enrichment	Employee count, industry, HQ — real figures, not parsed snippets
5	Business Scale Classification	Combines capital + headcount signals into MSME / SMB / Startup / Enterprise
6	Apollo + Lusha Contact Enrichment	Free search, 1 credit per email revealed
7	CRM Handoff	Raw facts only — no written hooks, no LLM-generated content

Everything in the pipeline is sourced free or near-free — no paid search engine dependency.

2. Step 1 — Discovery: MCA Names → Direct Domain Probe

No crt.sh, no Common Crawl, no SearXNG keyword search. One clean, free, official source. No search engine in the loop means no proxy dependency at any volume.

Step	Detail
MCA Company Master Data (data.gov.in)	Free, official — filtered to Active status, target NIC codes
Name → slug normalization	Strip "Pvt Ltd"/"LLP"; two slug variants per company (tight & hyphenated)
Direct HTTP probe	{slug}.com → .in → .co.in → ... — zero rate-limit or proxy exposure
Confidence check	Does the resolved page mention the company name? Rejects parked/squatted domains

3. Step 2 — Scrape: One Page Load, Multiple Jobs

Since Playwright is already loading the page for the consent check, the same pass captures everything else too — free.

Compliance Signals	Content & Tech Signals
Pre-consent cookies (analytics, ad pixels)	Company tagline (homepage hero)
Known CMP fingerprints (OneTrust, Cookiebot, etc.)	Careers page + open-role count
Banner + privacy policy presence	Payment gateway, CMS, analytics/marketing tools
Team page: founder/CEO name (internal cross-check)	Team headcount (named people on team page)

4. Step 3 — Scale Data: Apify LinkedIn Enrichment

Real employee counts, not parsed search snippets — and genuinely cheap.

Step	Detail
Resolved domain → LinkedIn company URL	Matched from company name / website field
Apify LinkedIn Company Scraper actor	No login, no cookies — reads public company-page data only
Returns	Employee count, industry, HQ, size band, website (structured JSON)
Cost	\$0.003–\$0.008 per company (third-party Apify pricing — verify live billing table before scaling)

5. Step 4 — Business Scale Classification & Pricing

Rule-based, free — no LLM needed. Combines MCA paid-up capital, Apify employee count, site team headcount, and NIC industry code into one bucket. This classification decides which plan gets pitched before the CRM person even opens the record.

Scale Bucket	Signal Band	Plan to Pitch	Price
MSME	< 20 employees, low paid-up capital	Free tier only	■0 — no sales push
SMB	20–50 employees, moderate capital	Starter	■60,000 / year
Startup / Growth	51–300 employees, confirmed revenue signal	Growth	~■5,00,000 ACV
Enterprise	100+ employees, BFSI / Healthcare, high revenue	Enterprise + SDF module	Custom, sales-assisted

BFSI / Healthcare note: these are "qualified enquiry" targets per the RFP, not self-serve signups — scale here confirms sales-worthiness, it doesn't pick a self-serve tier.

6. Step 5 — Contact Enrichment: Apollo + Lusha

Step	Detail
Flagged domain	From the consent-checker
Apollo api_search (free)	Filter: person_titles + domain
Filter has_email:true	Skip low-probability matches
people/match, reveal:true	1 credit — only on the chosen match

Title-match priority: 1. Founder/Co-founder → 2. CEO → 3. CTO → 4. Head of Engineering/Product → 5. Marketing Head

Site-scraped founder name overrides Apollo's title data on conflict — confirmed unreliable for smaller companies in live testing. Cost: 1 credit per record revealed.

7. Step 6 — What Reaches the CRM

Raw facts only. No written hooks, no LLM-generated content — the outreach writer builds the angle themselves. 11 fields, all outward-facing:

- company_name
- domain
- business_scale
- contact_name
- contact_title
- contact_email
- contact_source
- title_match_confidence
- consent_gap_summary
- company_tagline
- notable_open_roles

Everything else — capital, NIC code, tech stack, LinkedIn raw data — stays internal for ICP filtering only.

8. Throughput — Time & Cost to 1,000 Real Contacts

No shared search engine left in the pipeline, so the real bottleneck is now the Playwright scrape — fully under internal control.

Stage	Limited By	Throughput
MCA load + filter	Local compute only	Instant
Domain resolution (probe)	Your bandwidth / threads	~3,000–5,000 / hour
Playwright multi-purpose scrape	Local CPU / RAM, concurrent contexts	~800–1,200 / hour — bottleneck
Apify LinkedIn enrichment	Runs only on flagged candidates	Not the bottleneck at this volume
Apollo + Lusha enrichment	Runs only on confirmed matches	Untested rate limit, small volume

Conversion funnel: 1,000 MCA companies → 400 resolved → 200 flagged → 150 contactable (~15% overall conversion). To reach 1,000 real contacts: ~6,700 MCA companies fed in, ~2–3 hours of processing, \$15–\$34 total in Apify + Apollo/Lusha spend.

Conversion ratios are modeled estimates, not yet measured against the current MCA + direct-probe method — the 40% resolution rate in particular needs a real test batch to confirm.

9. Tooling & Cost Summary

Pipeline Stage	Tool	Cost
Discovery	MCA direct-probe (no search engine)	Free
Consent + content scrape	Playwright	Free (compute only)
Scale data	Apify LinkedIn Company Scraper	\$0.003–\$0.008 / company

Contact enrichment	Apollo, Lusha fallback	~\$1 / email credit *
--------------------	------------------------	-----------------------

** Search itself is free (0 credits, verified live). Reveal cost assumes paid-tier access — confirmed pending an actual test of the reveal endpoint. Total: \$15–\$34 to reach 1,000 real contacts.*

10. Next Steps

1. Re-pull MCA data filtered to the RFP's four target NIC codes (SaaS/IT, D2C, BFSI, Healthcare) — current shortlist is manufacturing-heavy
2. Verify Apollo's people/match reveal endpoint works on the current plan before relying on the cost model
3. Wire the Apify LinkedIn Company Scraper into the pipeline, verify its live billing table before scaling volume
4. Build the rule-based scale classifier from the four combined signals
5. Run a 100-company test batch end to end and measure real conversion rates at each stage

Summary generated from: consiva_pipeline_v2.pptx (Swaran Soft, Internal Build Reference v2).